

# assignment\_3\_solutions

January 11, 2023

## 1 Assignment 3: SQL and Pandas

### 1.1 Overview

Like Assignment 2, we will be using the Chinook database, this time focusing more on the sales side of the music store.

#### 1.1.1 Question 0

Set up your connection and cursor for `sqlite3`.

```
[ ]: import sqlite3
import pandas as pd
from pandasql import sqldf

connection = sqlite3.connect("../data/chinook.sqlite")
cursor = connection.cursor()
```

#### 1.1.2 Question 1

Using `pandas`' built-in function `read_sql_query(query, connection)`, query the entire `Customer` and `Invoice` tables, merging by `CustomerId`.

```
[ ]: customer_invoice_df = pd.read_sql_query('''
SELECT * FROM Customer
INNER JOIN Invoice
    ON Customer.CustomerId = Invoice.CustomerId
''', connection)

customer_invoice_df.head()
```

```
[ ]: CustomerId FirstName LastName Company Address City \
0          2    Leonie   Köhler    None Theodor-Heuss-Straße 34 Stuttgart
1          4    Bjørn   Hansen    None Ullevålsveien 14      Oslo
2          8     Daan  Peeters    None Grétrystraat 63     Brussels
3         14     Mark  Philips    Telus 8210 111 ST NW    Edmonton
4         23     John  Gordon    None 69 Salem Street   Boston

State Country PostalCode Phone ... SupportRepId InvoiceId \
```

0	None	Germany	70174	+49 0711 2842222	...	5	1
1	None	Norway	0171	+47 22 44 22 22	...	4	2
2	None	Belgium	1000	+32 02 219 03 03	...	4	3
3	AB	Canada	T6G 2C7	+1 (780) 434-4554	...	5	4
4	MA	USA	2113	+1 (617) 522-1333	...	4	5

	CustomerId	InvoiceDate	BillingAddress	BillingCity	\
0	2	2009-01-01 00:00:00	Theodor-Heuss-Straße 34	Stuttgart	
1	4	2009-01-02 00:00:00	Ullevålsveien 14	Oslo	
2	8	2009-01-03 00:00:00	Grêtrystraat 63	Brussels	
3	14	2009-01-06 00:00:00	8210 111 ST NW	Edmonton	
4	23	2009-01-11 00:00:00	69 Salem Street	Boston	

	BillingState	BillingCountry	BillingPostalCode	Total
0	None	Germany	70174	1.98
1	None	Norway	0171	3.96
2	None	Belgium	1000	5.94
3	AB	Canada	T6G 2C7	8.91
4	MA	USA	2113	13.86

[5 rows x 22 columns]

### 1.1.3 Question 2

Using `sqldf`, query a new DataFrame containing only the `CustomerId`, `FirstName`, `LastName`, `InvoiceDate`, and `Total`. **Drop any invoices before 2010.** Remember, when using an existing pandas DataFrame, you're query FROM the variable name, not the tables from previously.

```
[ ]: table = sqldf(''
    SELECT CustomerId, FirstName, LastName, InvoiceDate, Total
    FROM customer_invoice_df
    WHERE InvoiceDate > 2010
'')
table
```

```
[ ]:      CustomerId  FirstName  LastName  InvoiceDate  Total
0           43  Isabelle  Mercier  2010-01-08 00:00:00    1.98
1           45  Ladislav  Kovács  2010-01-08 00:00:00    1.98
2           47    Lucas  Mancini  2010-01-09 00:00:00    3.96
3           51   Joakim  Johansson  2010-01-10 00:00:00    6.94
4           57    Luis  Rojas  2010-01-13 00:00:00   17.91
..          ...      ...      ...      ...      ...
324          25   Victor  Stevens  2013-12-05 00:00:00    3.96
325          29   Robert   Brown  2013-12-06 00:00:00    5.94
326          35  Madalena  Sampaio  2013-12-09 00:00:00    8.91
327          44    Terhi  Hämäläinen  2013-12-14 00:00:00   13.86
328          58    Manoj  Pareek  2013-12-22 00:00:00    1.99
```

[329 rows x 5 columns]

## 1.2 Question 3

We now have a simplified list of invoices containing only name, id, date, and total. Let's figure out how much every has spent since 2010! Recall that pandas has the `groupby()` function (documentation [here](#)). Using this, make a new DataFrame grouping by the `CustomerId`, giving the total/sum for each.

```
[ ]: table.groupby(["CustomerId"]).sum()
```

```
[ ]:
      CustomerId      Total
1              1      39.62
2              2      12.87
3              3      39.62
4              4      28.73
5              5      38.64
6              6      40.71
7              7      40.64
8              8      30.69
9              9      31.68
10             10      28.71
11             11      21.78
12             12      36.63
13             13      25.74
14             14      28.71
15             15      22.78
16             16      36.63
17             17      27.74
18             18      37.62
19             19      13.87
20             20      39.62
21             21      27.72
22             22      39.62
23             23      14.85
24             24      43.62
25             25      35.69
26             26      45.64
27             27      28.71
28             28      27.78
29             29      36.63
30             30      31.68
31             31      28.71
32             32      21.78
33             33      36.63
```

34	27.74
35	37.62
36	21.78
37	42.63
38	25.74
39	38.62
40	13.87
41	37.62
42	28.73
43	40.62
44	32.71
45	45.62
46	38.69
47	35.64
48	31.71
49	21.78
50	36.63
51	32.68
52	28.71
53	21.78
54	36.63
55	25.74
56	37.62
57	30.78
58	38.62
59	26.74

### 1.2.1 Question 4

Turns out, the pandas `groupby()` statement has an equivalent in SQL! Putting

`GROUP BY <Table Name>.<Column Name>`

at the end of a query plus the keyword `SUM()` around the column you're grouping by creates a very similar table to the one we made in pandas!

Use what we've learned to query a table with `FirstName`, `LastName`, and `SUM(Total)`. Do the results match what we did in pandas?

```
[ ]: table = sqldf('''
    SELECT FirstName, LastName, SUM(Total)
    FROM customer_invoice_df
    WHERE InvoiceDate > 2010
    GROUP BY CustomerId
''')

table
```

[ ]:	FirstName	LastName	SUM(Total)
0	Luís	Gonçalves	39.62
1	Leonie	Köhler	12.87
2	François	Tremblay	39.62
3	Bjørn	Hansen	28.73
4	František	Wichterlová	38.64
5	Helena	Holý	40.71
6	Astrid	Gruber	40.64
7	Daan	Peeters	30.69
8	Kara	Nielsen	31.68
9	Eduardo	Martins	28.71
10	Alexandre	Rocha	21.78
11	Roberto	Almeida	36.63
12	Fernanda	Ramos	25.74
13	Mark	Philips	28.71
14	Jennifer	Peterson	22.78
15	Frank	Harris	36.63
16	Jack	Smith	27.74
17	Michelle	Brooks	37.62
18	Tim	Goyer	13.87
19	Dan	Miller	39.62
20	Kathy	Chase	27.72
21	Heather	Leacock	39.62
22	John	Gordon	14.85
23	Frank	Ralston	43.62
24	Victor	Stevens	35.69
25	Richard	Cunningham	45.64
26	Patrick	Gray	28.71
27	Julia	Barnett	27.78
28	Robert	Brown	36.63
29	Edward	Francis	31.68
30	Martha	Silk	28.71
31	Aaron	Mitchell	21.78
32	Ellie	Sullivan	36.63
33	João	Fernandes	27.74
34	Madalena	Sampaio	37.62
35	Hannah	Schneider	21.78
36	Fynn	Zimmermann	42.63
37	Niklas	Schröder	25.74
38	Camille	Bernard	38.62
39	Dominique	Lefebvre	13.87
40	Marc	Dubois	37.62
41	Wyatt	Girard	28.73
42	Isabelle	Mercier	40.62
43	Terhi	Hämäläinen	32.71
44	Ladislav	Kovács	45.62
45	Hugh	O'Reilly	38.69

46	Lucas	Mancini	35.64
47	Johannes	Van der Berg	31.71
48	Stanisław	Wójcik	21.78
49	Enrique	Muñoz	36.63
50	Joakim	Johansson	32.68
51	Emma	Jones	28.71
52	Phil	Hughes	21.78
53	Steve	Murray	36.63
54	Mark	Taylor	25.74
55	Diego	Gutiérrez	37.62
56	Luis	Rojas	30.78
57	Manoj	Pareek	38.62
58	Puja	Srivastava	26.74

[ ]: